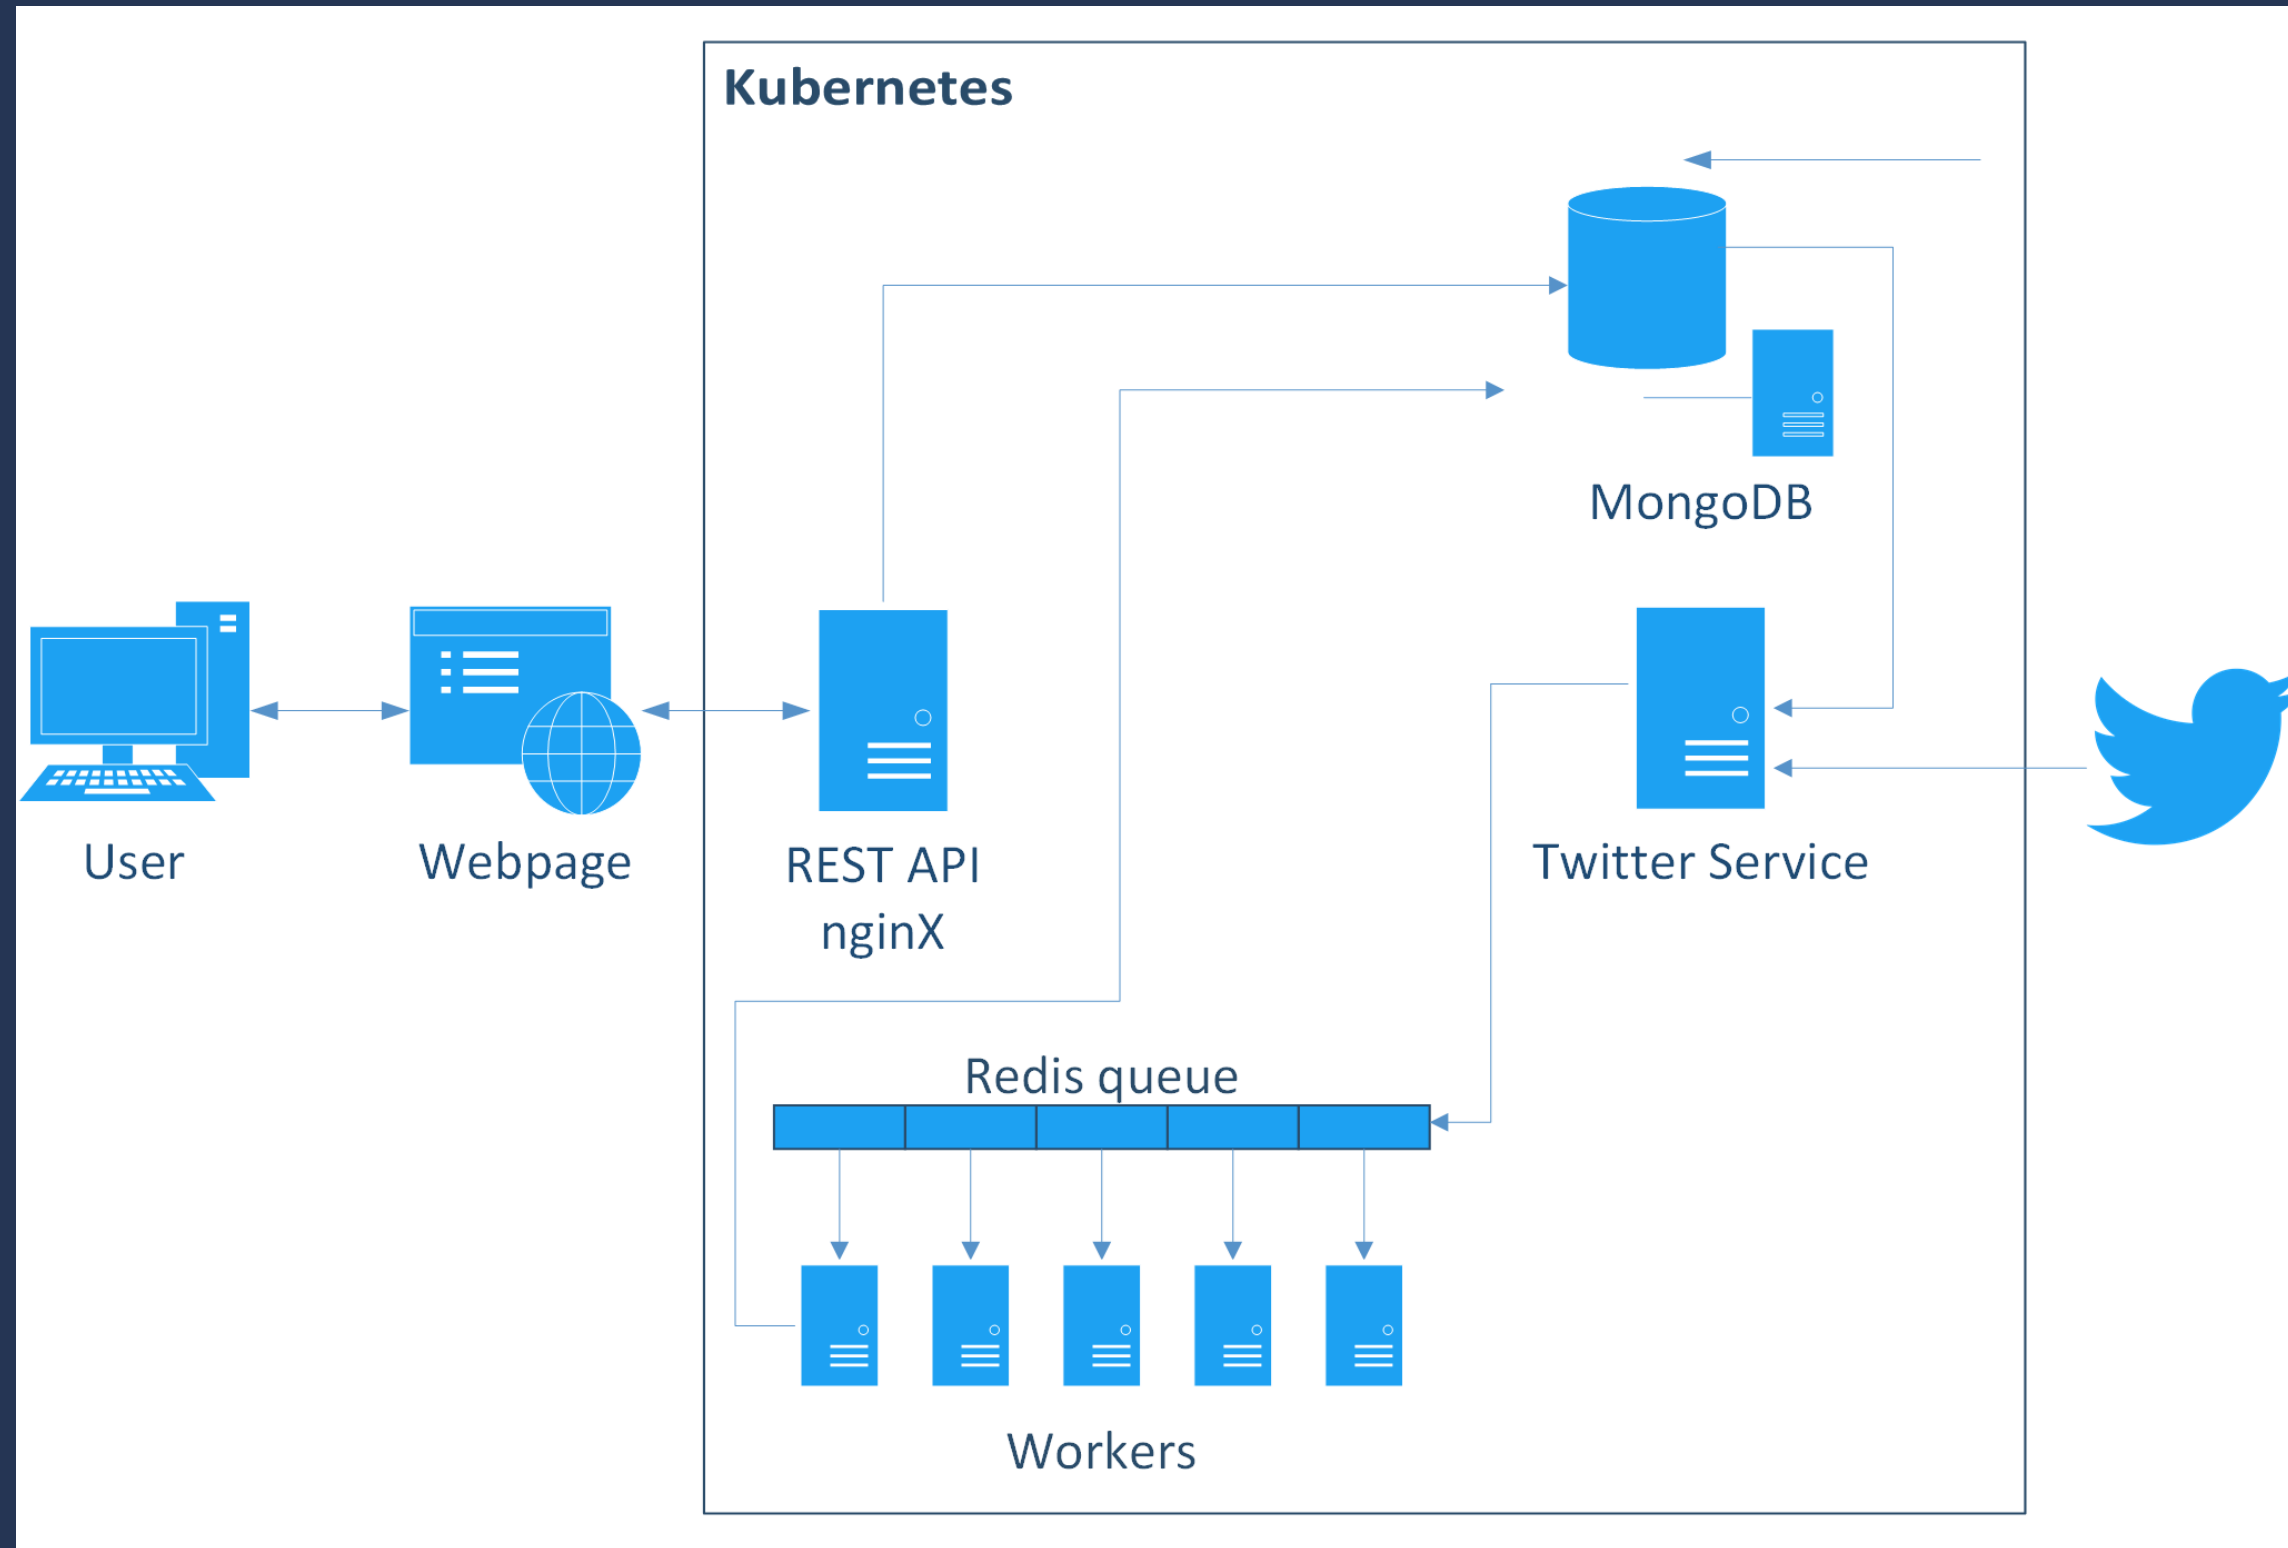


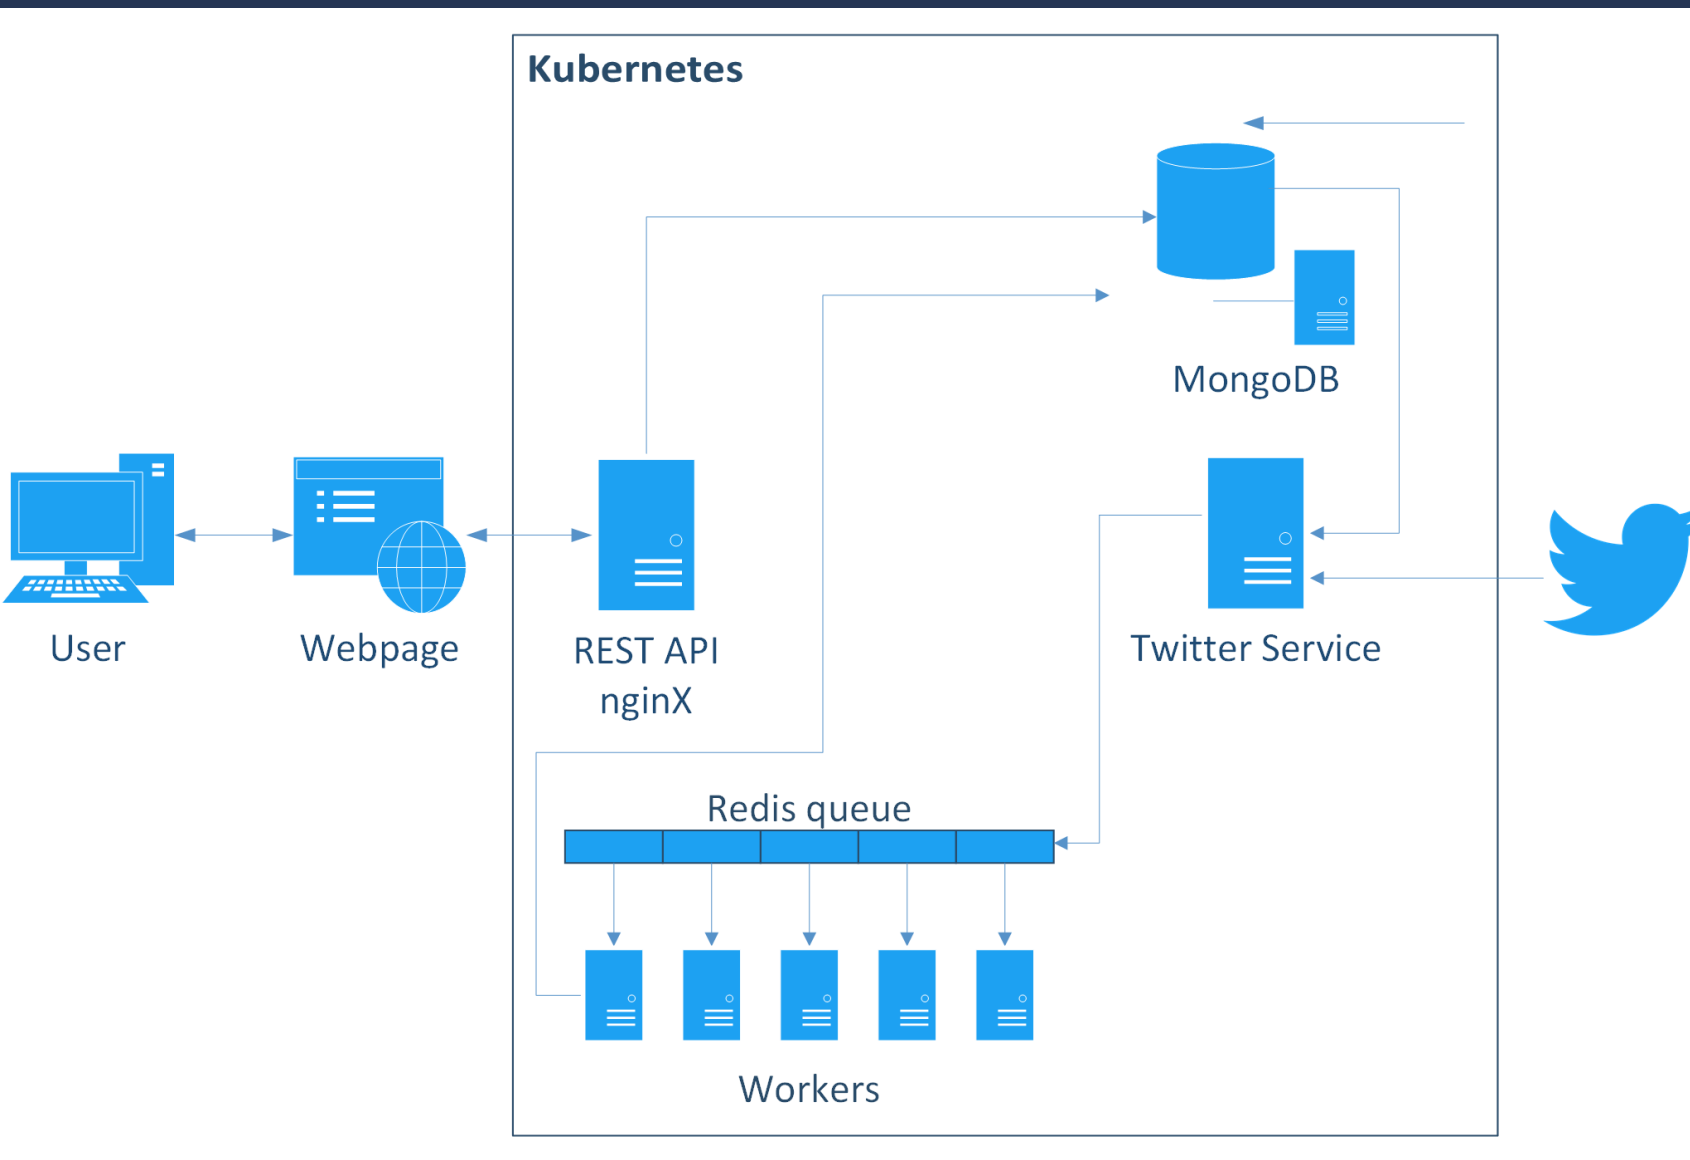
Project Presentation Stage 1

**Stefanie Ziltener, Marc Heimgartner, Benjamin
Bürgisser, Simon Tännler**

Advanced Software Engineering FS 2017, University of Zürich

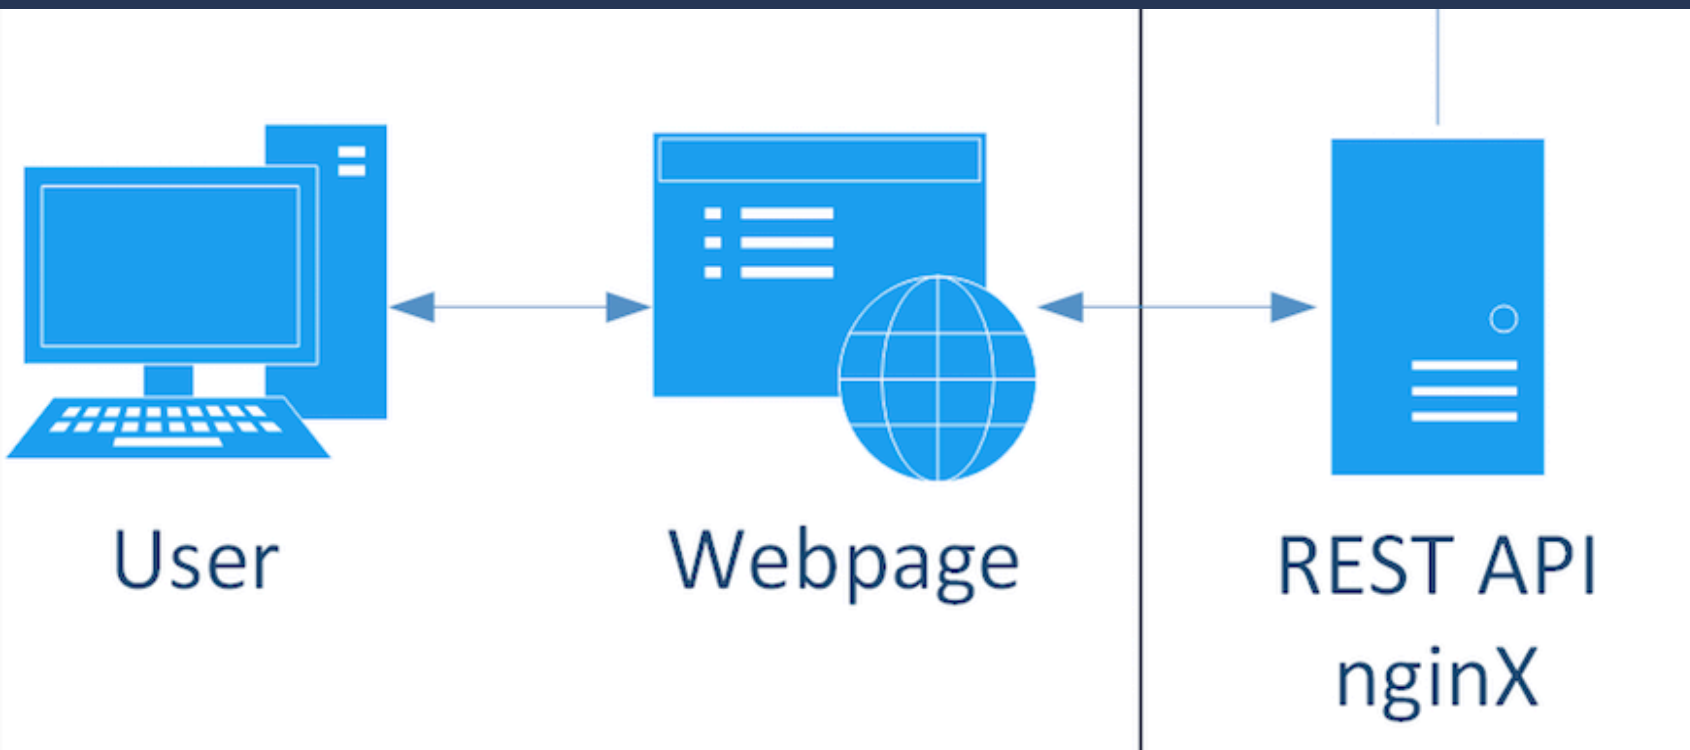


Containerized Microservices deployed through Kubernetes




- Components embedded in (Docker) Containers
- Containers have (ideally) one single responsibility

Frontend and API



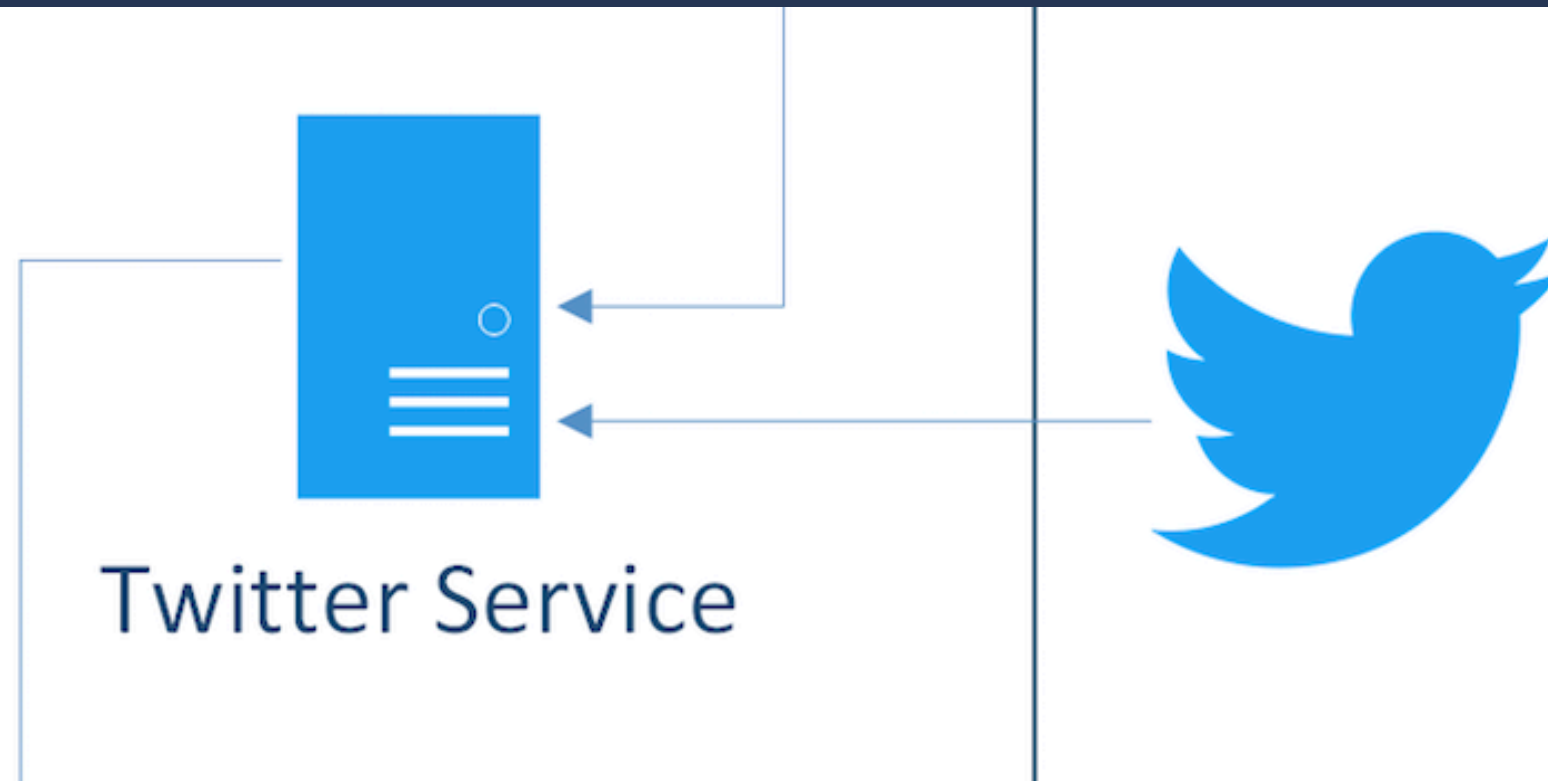
- User inputs term
- Click on register
- Send request to API

Timeseries: MongoDB (I)




- Request Handler stores Term in MongoDB
- Persistency guaranteed by GCE Persistence Disk

Twitter Service




- Gets notified of newly created Terms
- Streaming stops and restarts with the new Term added for tracking
- Arriving Tweets are immediately stored into the Queue

Worker Queue



- Redis: in-memory data structure store
- A FIFO queue of Strings (Tweets)
- Load generator
 - API endpoint to add Strings to Queue directly

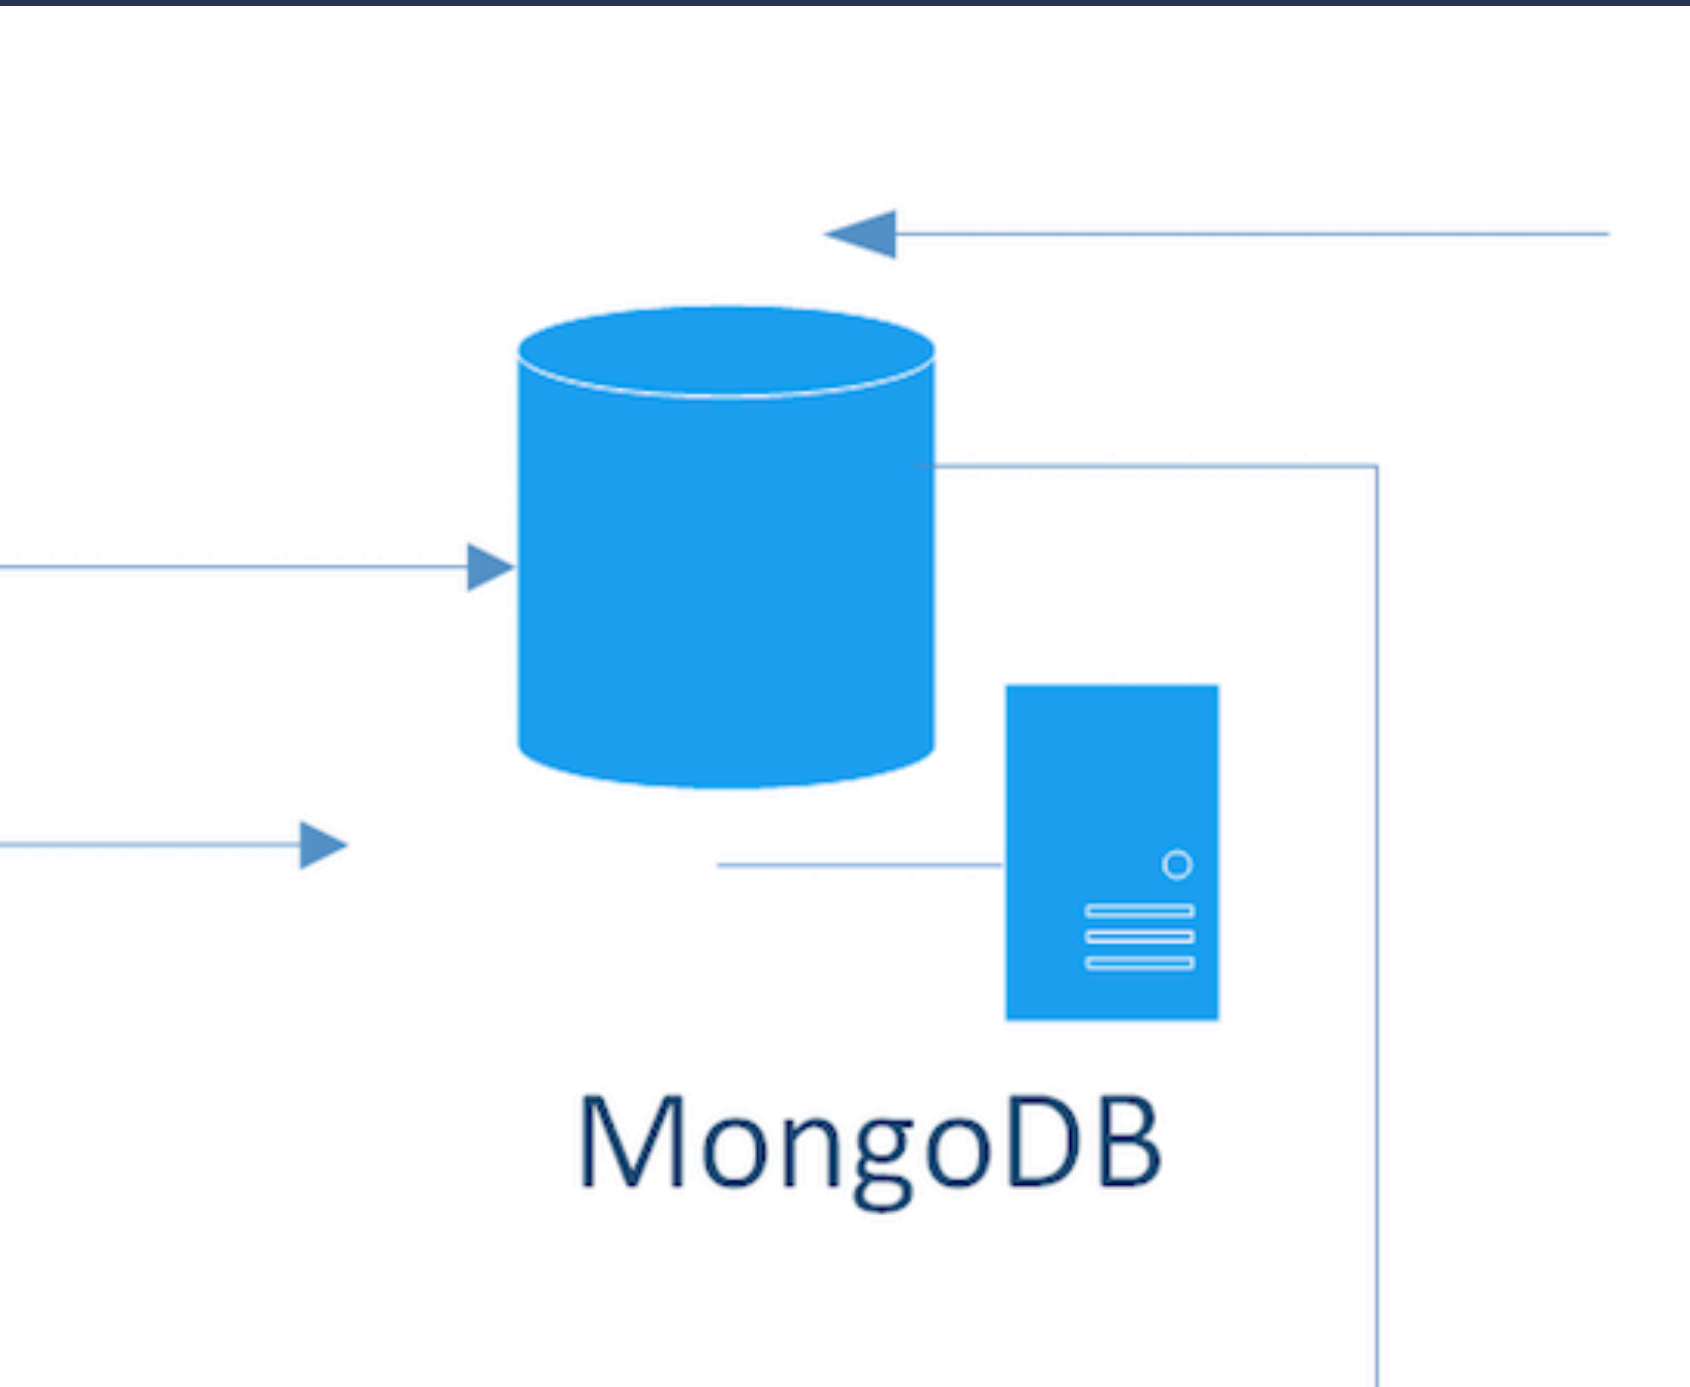
Compute Workers



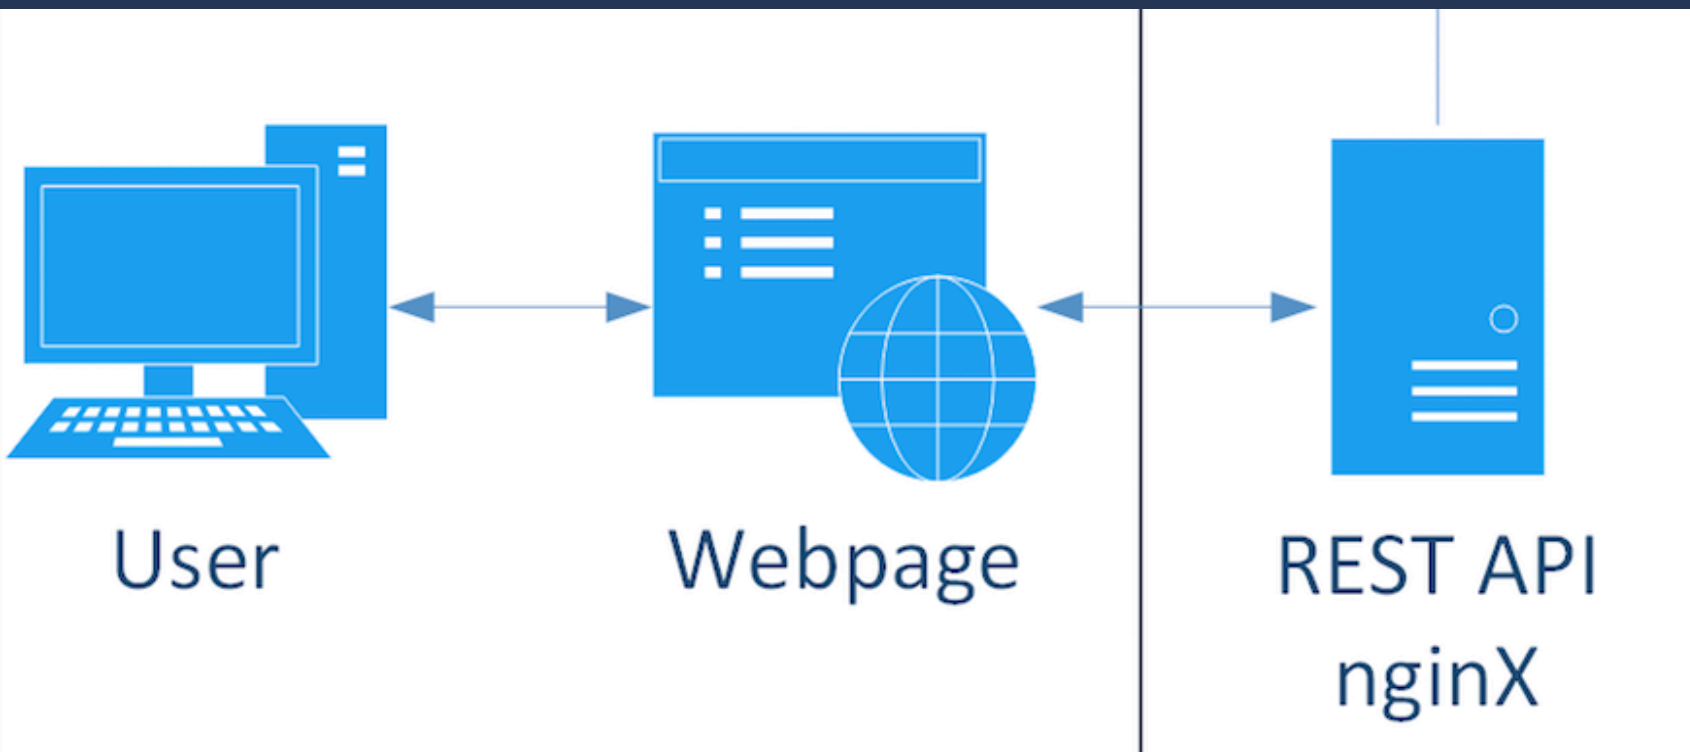
- Running Workers process the Queue:
 - Assign Tweet to Term (filtering)
 - Calculate Sentiment
- Length of the Queue defined the number of Workers
 - Scaled through Kubernetes.

Timeseries: MongoDB (II)

- Workers store the calculated Sentiment into MongoDB.



Displaying Results



- Rest API gets request
- Collect relevant data from MongoDB
- Browser renders data
- Socket gets opened for continues pushes


Architectural Styles

- Client / Server through Rest API
- Event-Driven notification of Term updates
- Blackboard: Redis Queue
 - Factory: Twitter Service
 - Worker: Compute Workers
- Highly decoupled Microservices

Do you even scale?

Initial Idea

- Digital Ocean
- Docker Swarm



docker

What does Kubernetes do for us

- Every component is potentially scalable through Kubernetes
 - Even MongoDB!
- Fault Tolerance:
 - Container recovery through Kubernetes
 - Decoupled design and Microservice
- Elasticity
 - Container scaling through Kubernetes

Concerns

- High lock-in to Kubernetes
- Twitter: 400 keyword tracking per stream
 - "default access level allows up to 400 track keywords"
- Matching: may not match to terms perfectly
- Redis: may become a bottleneck
 - but we highly doubt it

Technology Zoo - Platform

- Cloud Platform: Google Container Engine (GCE)
 - Easy support of Kubernetes
- Containerization: Docker
 - Popular Container engine
- Container orchestration: Kubernetes
 - Popular Container orchestration

Technology Zoo - Backend

- Programming Language: Google Go
 - New Programming language 🎉
 - Uniquely suited for Web development
 - Have I mentioned it's fast?
- Terms Storage: MongoDB
 - Easy data schema
- Queue Storage: Redis

Technology Zoo - Frontend

- Frontend: Vuejs
 - Similar to Angular and React
- Webserver: nginx
 - Battle-proven Webserver

Dev Environment

Docker containers using Docker Compose

Demo

Questions?